

## K - Nearest Neighbors (K-NN)

**Aim:**

1. Implement the K-NN algorithm for classification or regression.
2. Apply the K-NN algorithm to a given dataset and predict the class or value for test data.
3. Evaluate the accuracy or error of the predictions and analyze the results

<b>Theory:</b>	<p>K-Nearest Neighbors (K-NN) is a simple but powerful machine learning algorithm used for both classification and regression tasks.</p> <ol style="list-style-type: none"> <li>1. Intuition: K-NN is based on the idea that objects (data points) that are close to each other in a feature space are more likely to belong to the same class or have similar values (for regression).</li> <li>2. How it works: <ul style="list-style-type: none"> <li>- Classification: Given a new data point, K-NN finds the K nearest data points in the training dataset and assigns the class label that is most common among those K neighbors to the new point.</li> <li>- Regression: For regression tasks, K-NN calculates the average (or another aggregation) of the target values of the K nearest neighbors and assigns this value to the new data point.</li> </ul> </li> <li>3. Hyperparameter K: The choice of the hyperparameter K (the number of neighbors to consider) is critical. A small K may lead to a noisy model (sensitive to outliers), while a large K may lead to a biased model (smoothing over variations in the data). K is typically an odd number to avoid ties in voting.</li> <li>4. Distance Metric: K-NN uses a distance metric (e.g., Euclidean distance, Manhattan distance, etc.) to measure the similarity between data points. The choice of distance metric should be appropriate for your data and problem.</li> <li>5. Scaling Features: It's important to scale features before applying K-NN, especially when using distance-based metrics, to ensure that all features have equal influence on the results.</li> <li>6. Pros: <ul style="list-style-type: none"> <li>- Simple and easy to understand.</li> <li>- No assumptions about the data distribution.</li> <li>- Works well for both classification and regression tasks.</li> <li>- Non-parametric (does not make assumptions about the functional form of relationships).</li> </ul> </li> <li>7. Cons: <ul style="list-style-type: none"> <li>- Can be computationally expensive, especially for large datasets.</li> <li>- Sensitive to the choice of K and the distance metric.</li> <li>- Requires a sufficient amount of training data.</li> <li>- May not perform well when the feature space is high-dimensional.</li> </ul> </li> <li>8. Use Cases:</li> </ol>
----------------	--

	<ul style="list-style-type: none"> <li>- K-NN is often used for tasks such as recommendation systems, image classification, and anomaly detection.</li> <li>- It can be used as a baseline model for comparison with more complex algorithms.</li> </ul> <p>9. Model Evaluation: Common evaluation metrics for K-NN include accuracy (for classification) and mean squared error (for regression). Cross-validation is often used to estimate the model's generalization performance.</p> <p>10. Implementation: Python libraries like scikit-learn provide easy-to-use implementations of K-NN for both classification and regression.</p> <p>In summary, K-NN is a versatile and interpretable algorithm suitable for various tasks. It's important to choose appropriate values for K and the distance metric based on your data and problem domain to achieve good performance.</p>
--	---

<b>Dataset:</b>	<p>We use the iris dataset which is available in the public domain, although we can also create our own datasets</p> <p>We download the iris dataset and save in .csv format</p>
-----------------	--

The following is the Python code to implement K-NN Classifier

(To run this code, Scikit-Learn must be installed (pip install scikit-learn) and Pandas must be installed (pip install pandas).

<b>Python Code</b>	
--------------------	--

<b>Output</b>	
---------------	--

For Video demonstration of the practical click on the link or scan the QR-code



## PRACTICAL NO: 8

**Aim:- Implement the K-NN Algorithm for classification or regression.**

**Apply K-NN Algorithm on the given dataset & predict the class or value for test data.**

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

```
In [2]: df = pd.read_csv('C:/Users/RDNC/Desktop/diabetes.csv')
df.head()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

```
In [3]: df.shape
```

```
Out[3]: (768, 9)
```

```
In [4]: df.dtypes
```

```
Out[4]: Pregnancies      int64
Glucose      int64
BloodPressure  int64
SkinThickness int64
Insulin      int64
BMI          float64
DiabetesPedigreeFunction float64
Age          int64
Outcome      int64
dtype: object
```

```
In [9]: x = df.drop('Outcome',axis=1).values
y = df['Outcome'].values
```

```
In [18]: from sklearn.model_selection import train_test_split
```

```
In [19]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.4,random_state=42, s1
```

```
In [24]: from sklearn.neighbors import KNeighborsClassifier
neighbors = np.arange(1,9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

for i,k in enumerate(neighbors):
    #Setup a knn classifier with k neighbors
    knn = KNeighborsClassifier(n_neighbors=k)
    #Fit the model
    knn.fit(X_train, y_train)
    #Compute accuracy on the training set
```



```

train_accuracy[i] = knn.score(X_train, y_train)
#Compute accuracy on the test set
test_accuracy[i] = knn.score(X_test, y_test)

```


```

In [25]: plt.title('k-NN Varying number of neighbors')
plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.plot(neighbors, train_accuracy, label='Training accuracy')
plt.legend()
plt.xlabel('Number of neighbors')
plt.ylabel('Accuracy')
plt.show()

```

File:///C:/Users/Lenovo/Downloads/Practical no 8- K-NEAREST NEIGHBOUR.html

4/36



```
In [27]: knn = KNeighborsClassifier(n_neighbors=7)
```

```
In [30]: knn.fit(X_train,y_train)
```

```
Out[30]: KNeighborsClassifier(n_neighbors=7)
```

```
In [32]: knn.score(X_test,y_test)
```

C:\Users\RDNC\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. 'skew', 'kurtosis'), the default behavior of 'mode' typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of 'keepdims' will become False, the 'axis' over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set 'keepdims' to True or False to avoid this warning.

```
mode, _ = stats.mode(y[neigh_ind, k], axis=1)
```

```
Out[32]: 0.7142857142857143
```

```
In [36]: from sklearn.metrics import confusion_matrix
```

```
In [ ]:
```

```
In [37]: y_pred = knn.predict(X_test)
```

```
In [39]: confusion_matrix(y_test,y_pred)
```

```
Out[39]: array([[163,  43],  
               [ 45,  57]], dtype=int64)
```

```
In [43]: from sklearn.metrics import classification_report
```

```
In [ ]:
```

```
In [44]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.79	0.79	206
1	0.57	0.56	0.56	102
accuracy			0.71	308
macro avg	0.68	0.68	0.68	308
weighted avg	0.71	0.71	0.71	308


```
In [46]: y_pred_proba = knn.predict_proba(X_test)[:,-1]
```

```
In [48]: from sklearn.metrics import roc_curve
```

```
In [50]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
```

```
In [52]: plt.plot([0,1],[0,1], 'k--')  
plt.plot(fpr,tpr, label='Knn')  
plt.xlabel('fpr')  
plt.ylabel('tpr')  
plt.title('Knn(n_neighbors=7) ROC curve')  
plt.show()
```

Knn(n\_neighbors=7) ROC curve



```
In [54]: from sklearn.metrics import roc_auc_score
roc_auc_score(y_test,y_pred_proba)
```

```
Out[54]: 0.7536645726251665
```

```
In [56]: from sklearn.model_selection import GridSearchCV
```

```
In [58]: param_grid = {'n_neighbors':np.arange(1,50)}
```

```
In [61]: knn = KNeighborsClassifier()
knn_cv= GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(x,y)
```

```
1]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
      param_grid={'n_neighbors': array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
      18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,
      35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49])})
```

```
3]: knn_cv.best_score_
```

```
3]: 0.7578558696205755
```

```
4]: knn_cv.best_params_
```

```
4]: {'n_neighbors': 14}
```